




中国科学院大学  
University of Chinese Academy of Sciences

# The TAU Profiler

## Part1: Introduction



人工智能技术学院

缪青海

miaoqh@ucas.ac.cn



# TAU

Home  
About  
News  
Documentation  
Wiki  
Downloads  
Publications  
Contact  
Applications  
CCA  
KTAU  
External links  
Search

P•R•L  
Performance Research Lab

Tuning and Analysis Utilities

TAU Performance System® is a portable profiling and tracing toolkit for performance analysis of parallel programs written in Fortran, C, C++, UPC, Java, Python.

TAU (Tuning and Analysis Utilities) is capable of gathering performance information through instrumentation of functions, methods, basic blocks, and statements as well as event-based sampling. All C++ language features are supported including templates and namespaces. The API also provides selection of profiling groups for organizing and controlling instrumentation. The instrumentation can be inserted in the source code using an automatic instrumentor tool based on the Program Database Toolkit (PDT), dynamically using DyninstAPI, at runtime in the Java Virtual Machine, or manually using the instrumentation API.


TAU's profile visualization tool, paraprof, provides graphical displays of all the performance analysis results, in aggregate and single node/context/thread forms. The user can quickly identify sources of performance bottlenecks in the application using the graphical interface. In addition, TAU can generate event traces that can be displayed with the Vampir, Paraver or JumpShot trace visualization tools.

- <https://www.cs.uoregon.edu/research/tau/home.php>



# Content

- Overview
- Instrumentation
- Profiling
- Tracing
- Analyzing
- Application Scenario



# TAU Performance System<sup>®</sup>

- **T**uning and **A**nalysis **U**tilities (TAU)
- Comprehensive performance profiling and tracing
  - ☐ Integrated, scalable, flexible, portable
  - ☐ Targets all parallel programming/execution paradigms
- Integrated performance toolkit
  - ☐ Instrumentation, measurement, analysis, visualization
  - ☐ Widely-ported performance profiling / tracing system
  - ☐ Performance data management and data mining
  - ☐ Open source (BSD-style license)
- Easy to integrate in application frameworks



# What is TAU?

- TAU is a performance evaluation tool. It supports parallel **profiling and tracing**
  - **Profiling** shows you how much (total) time was spent in each routine;
  - **Tracing** shows you *when* the events take place in each process along a timeline.
- Profiling and tracing can measure time as well as hardware performance counters (cache misses, instructions) from your CPU



# What is TAU?

- TAU can automatically **instrument** your source code using a package called PDT for routines, loops, I/O, memory, phases, etc.
- TAU runs on most HPC platforms and it is free (BSD style license)
- TAU has instrumentation, measurement and analysis tools
  - paraprof is TAU' s 3D profile browser
- To use TAU' s automatic source instrumentation, you may set a couple of environment variables and substitute the name of your compiler with a TAU shell script



# What Can TAU Do?

- How much time is spent in each application routine and outer *loops*? Within loops, what is the contribution of each *statement*?
- How many instructions are executed in these code regions? Floating point, Level 1 and 2 *data cache misses*, hits, branches taken?
- What is the *peak heap memory* usage of the code? When and where is memory allocated/de-allocated? Are there any memory leaks?



# What Can TAU Do?

- How much time does the application spend performing I/O? What is the peak read and write *bandwidth* of individual calls, total volume?
- What is the contribution of different *phases* of the program? What is the time wasted/spent waiting for collectives, and I/O operations in Initialization, Computation, I/O phases?
- How does the application *scale*? What is the efficiency, runtime breakdown of performance across different core counts?





# Some Features in TAU

- *tau\_exec*: A tool to simplify TAU's usage
- PDT: New parsers from ROSE Compiler, LLNL [Dan Quinlan, CASC]
- Power profiling in TAU
- Support for accelerators: CUDA, OpenCL, Intel Xeon Phi Co-processor
- Event based sampling
- OpenMP instrumentation using OpenMP Tools Interface with Intel compilers (OMPT)
- Support for Intel TBB, CILK, and MPC [[paratools.com/mpc](http://paratools.com/mpc)]
- ParaProf 3D topology displays



# What does TAU support?

**C/C++** **Python** **OpenCL**  
**CUDA** **UPC**  
**Fortran** **OpenACC** **GPI** **MPI**  
**pthread** **Java**  
**Intel MIC** **OpenMP**  
**Intel** **GNU** **PGI** **Cray** **Sun**  
**MinGW** **LLVM** **Windows** **AIX**  
**OS X** **Linux**  
**NVIDIA Kepler** **ARM**



# How does TAU work?

## Instrumentation:

- Adds probes to perform measurements
  - Source code instrumentation using pre-processors and compiler scripts
  - Wrapping external libraries (I/O, MPI, Memory, CUDA, OpenCL, pthread)
  - Rewriting the binary executable



# How does TAU work?

## Measurement:

- Profiling or Tracing using **wallclock time** or **hardware counters**:
  - Direct instrumentation (Interval events measure exclusive or inclusive duration)
  - Indirect instrumentation (Sampling measures statement level contribution)
  - Throttling and runtime control of low-level events that execute frequently
  - Per-thread storage of performance data
  - Interface with external packages (e.g. PAPI hw performance counter library)

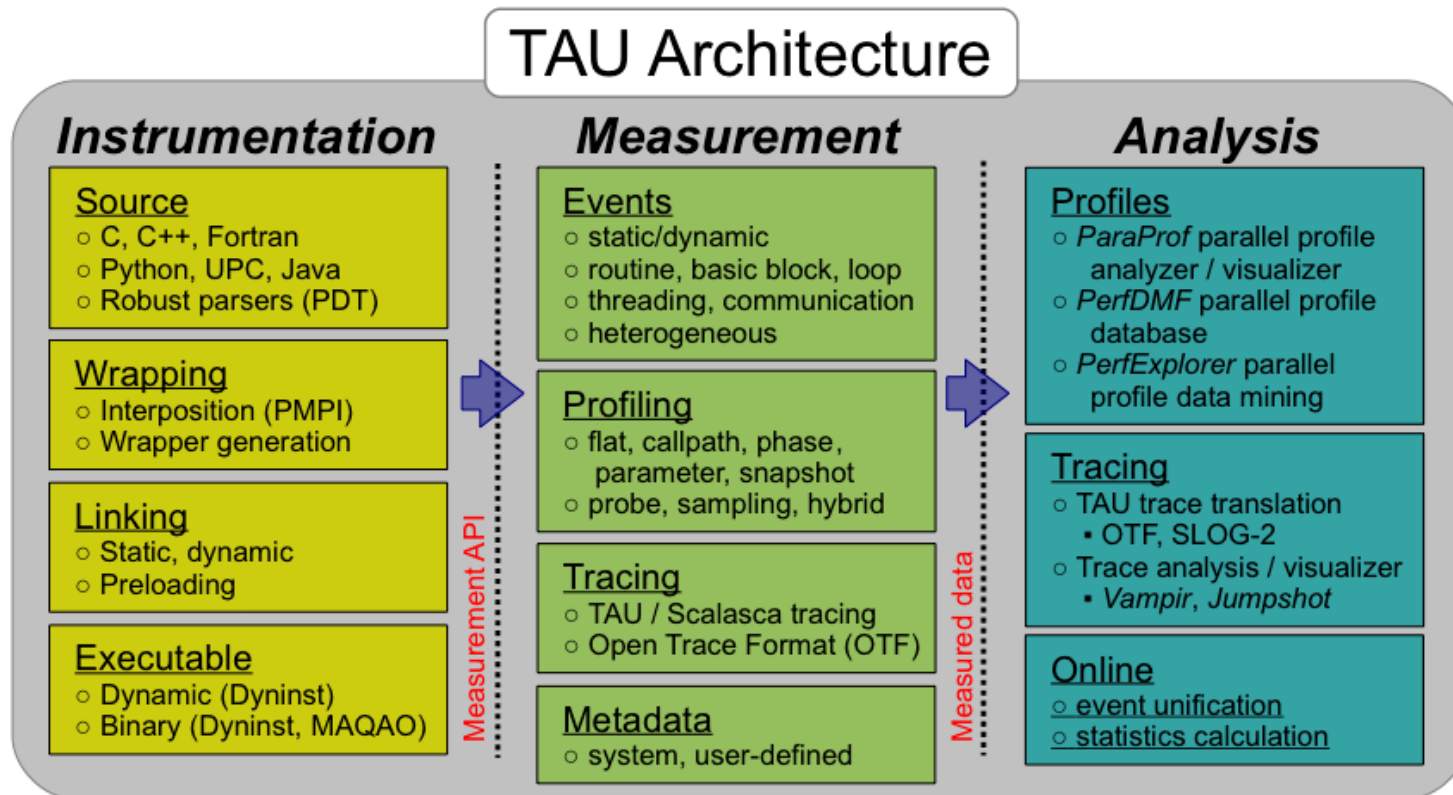


# How does TAU work?

## Analysis:

- Visualization of profiles and traces
  - 3D visualization of profile data in paraprof, perfexplorer tools
  - Trace conversion & display in external visualizers (Vampir, Jumpshot, ParaVer)

# How does TAU work?







# Profiling and Tracing

## Profiling

Value: Exclusive  
Units: seconds



## Tracing

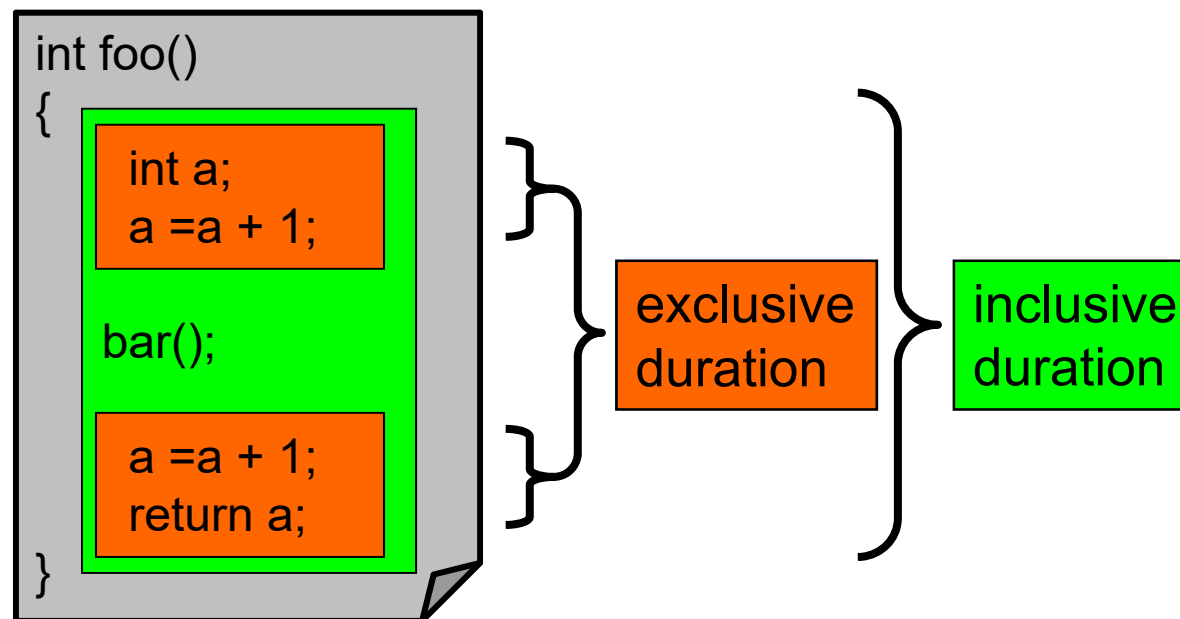


- **Profiling** shows you **how much** (total) time was spent in each routine
- Tracing shows you when the events take place on a timeline
- Metrics can be time or hardware performance counters (cache misses, instructions)
- TAU can automatically instrument your source code using a package called PDT for routines, loops, I/O, memory, phases, etc.



# Inclusive vs. Exclusive Measurements


- Performance with respect to code regions
- Exclusive measurements for region only
- Inclusive measurements includes child regions







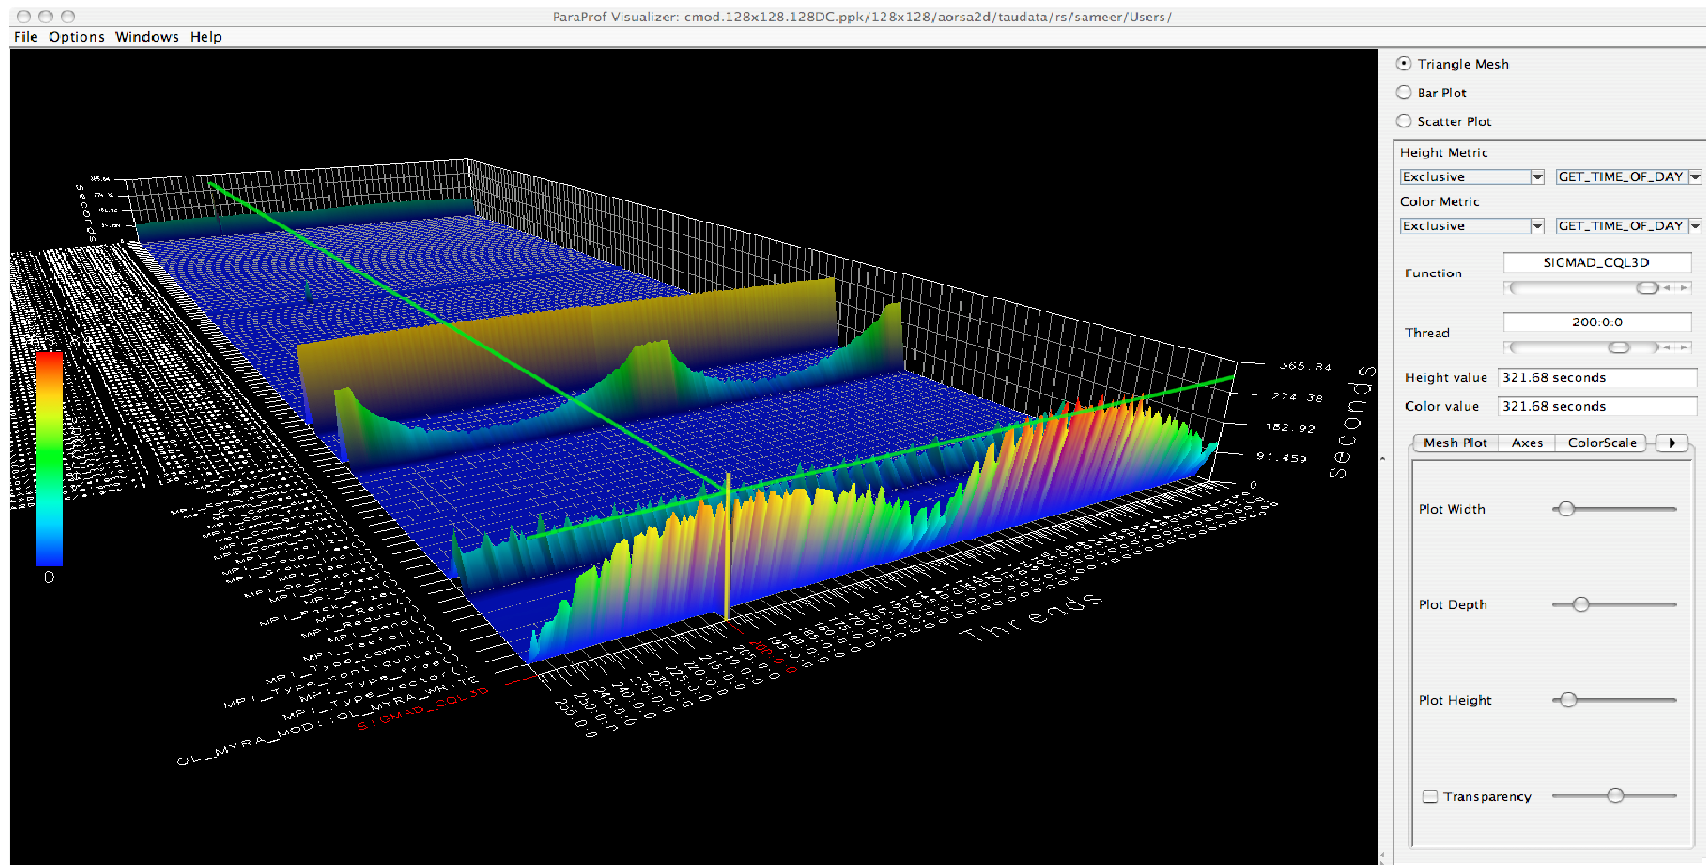
# Identifying Potential Bottlenecks



low mflops in loops?

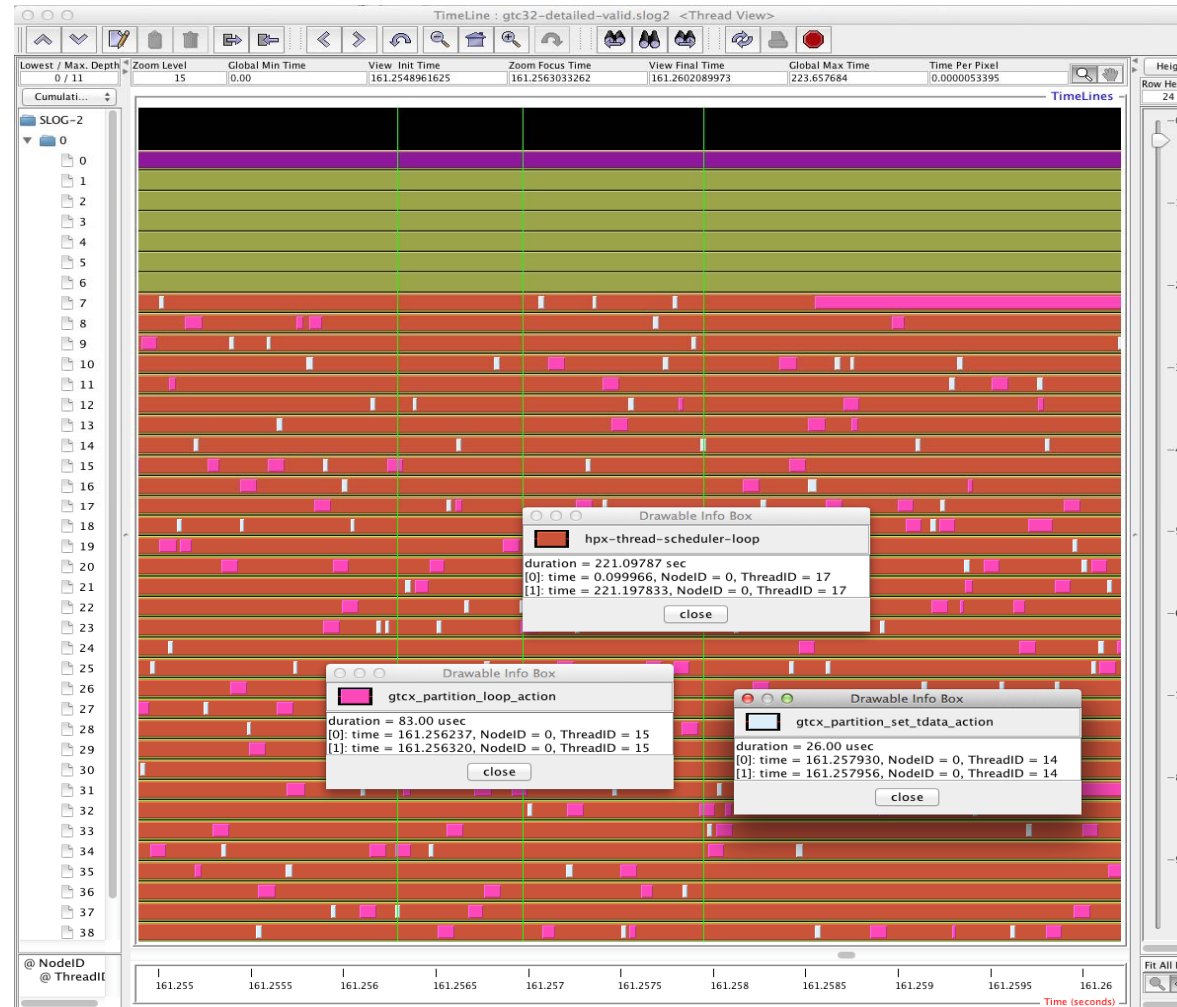


# ParaProf 3D Profile Browser






# Jumpshot Trace Visualizer in TAU





# Tracing Communication in Jumpshot



```
% export TAU_MAKEFILE=$TAU/Makefile.tau-icpc-papi-mpi-pdt
% cmake -DCMAKE_CXX_COMPILER=tau_cxx.sh; make -j 8
% export TAU_TRACE=1
% mpirun -np 16 ./a.out ; tau_treemerge.pl; tau2slog2 tau.trc tau.edf -o a.slog2
% jumpshot a.slog2 &
```



# Content

- Overview
- Instrumentation
- Profiling
- Tracing
- Analyzing
- Application Scenario



# Types of Instrumentation

- TAU provides three methods:
  - ☐ <1>Library interposition using tau\_exec.
  - ☐ <2>Compiler directives.
  - ☐ <2>Source transformation using PDT.

<i>Method</i>	<b>Requires recompiling</b>	<b>Requires PDT</b>	<b>Shows MPI events</b>	<b>Routine-level event</b>	<b>Low level events (loops, phases, etc...)</b>	<b>Throttling to reduce overhead</b>	<b>Ability to exclude file from instrumentation</b>
Interposition			Yes			Yes	
Compiler	Yes		Yes	Yes		Yes	Yes
Source	Yes	Yes	Yes	Yes	Yes	Yes	Yes



## <1> Using tau\_exec

- Dynamic instrumentation:
  - Achieved through library pre-loading.
  - Can be used on both uninstrumented and instrumented binaries.
  
- Place tau\_exec before the executable:
  - %> tau\_exec -io ./a.out
  - %> mpirun -np 4 tau\_exec -io ./a.out



# TAU scripted compilation

- Tau\_exec:
  - Simple but less detailed.
  
- For more detailed profiles, TAU provides two means to **compile** your application with TAU:
  - **Compiler Based** Instrumentation
    - through your compiler.
  - **Source Based** Instrumentation
    - through source transformation using PDT.





# Compiler Based Instrumentation

- TAU provides these scripts:
  - tau\_cc.sh, tau\_cxx.sh, tau\_upc.sh, tau\_f77.sh and tau\_f90.sh to compile programs.
- You might use tau\_cc.sh to compile a C program by typing:

```
%> module load tau
%> tau_cc.sh -tau_options=-optComplnst
        samplecprogram.c
```



# Compiler Based Instrumentation

- On machines where a TAU module is not available, you will need to set the tau makefile and/or options:

```
%>tau_cc.sh -tau_makefile=[path to makefile] \  
-tau_options=[option] samplecprogram.c
```

- Makefile:
  - ☐ can be found in the /[arch]/lib directory of your TAU distribution.
  - ☐ for example / x86\_64/lib/Makefile.tau-mpi-pdt .



# Compiler Based Instrumentation

- You can also use a Makefile specified in an environment variable.
- To run tau\_cc.sh so it uses the Makefile specified by environment variable TAU\_MAKEFILE.

```
%>export TAU_MAKEFILE=[path to tau]/[arch]/lib/[makefile]
```

```
%>export TAU_OPTIONS=-optCompInst
```

```
%>tau_cc.sh sampleCprogram.c
```



# Source Based Instrumentation

- TAU provides these scripts:
  - tau\_cc.sh, tau\_cxx.sh, tau\_upc.sh, tau\_f77.sh and tau\_f90.sh to compile programs.
- You might use tau\_cc.sh to compile a C program by typing:

```
%> module load tau  
%> tau_cc.sh samplecprogram.c
```

- NOTE:
  - **When setting the TAU\_MAKEFILE make sure the Makefile name contains pdt.**



# Options to TAU compiler scripts

- There are some commonly used options available to the TAU compiler scripts.
- Either set them via the **TAU\_OPTIONS** environment variable or the **-tau\_options = option** to tau\_cc.sh.
- -optVerbose
  - ☐ Enable verbose output (default: on)
- -optKeepFiles
  - ☐ Do not remove intermediate files
- -optShared
  - ☐ Use shared library of TAU (consider when using tau\_exec)



# Content

- Overview
- Instrumentation
- Profiling
- Tracing
- Analyzing
- Application Scenario



# Profiling

- Profiling shows the **summary statistics of performance** metrics that characterize application performance behavior.
- Examples of performance metrics are:
  - the CPU time associated with a routine,
  - the count of the secondary data cache misses associated with a group of statements,
  - the number of times a routine executes, etc.



# Running the Application

- After instrumentation and compilation are completed, the profiled application is run to generate the profile data file.
  - By default, profiles are placed in the current directory or set PROFILEDIR.
  - set the TAU\_VERBOSE to see the steps the TAU takes when your application is running.

```
% setenv TAU_VERBOSE 1  
% setenv PROFILEDIR /home/sameer/profiledata  
% mpirun -np 4 matrix
```





# Reducing Overhead

- TAU automatically throttles short running functions in an effort to reduce the amount of overhead.
  - This feature may be turned off by setting the environment variable `TAU_THROTTLE=0`.
- The default rules TAU uses to determine which functions to throttle is:
  - `numcalls > 100000 && usecs/call < 10`

```
% setenv TAU_THROTTLE_NUMCALLS 2000000  
% setenv TAU_THROTTLE_PERCALL 5
```



# Profiling each event callpath

- You can enable callpath profiling by setting the environment variable:

**TAU\_CALLPATH**

- In this mode TAU will record the each event callpath to the depth set by the environment variable (default is two).

**TAU\_CALLPATH\_DEPTH**

- Because instrumentation overhead will increase with the depth of the callpath, you should use the shortest call path that is sufficient.



# Using Hardware Counters

- Performance counters exist on many modern microprocessors.
  - cache misses, floating point operations, etc.
- The **Performance Data Standard and API (PAPI)** package provides a uniform interface to access these performance counters.[<http://icl.cs.utk.edu/papi/>].
- To use these counters, you must:
  - First, find out which PAPI events your system supports.
  - Next, test the compatibility between each metric you wish papi to profile.
  - Next, make sure that you are using a makefile with papi in its name.



# Content

- Overview
- Instrumentation
- Profiling
- Tracing
- Analyzing
- Application Scenario



# Tracing

- Typically, profiling shows the distribution of execution time across routines.
  - It can show the code locations associated with specific bottlenecks,
  - but it can not show the temporal aspect of performance variations.
- Tracing the execution of a parallel program shows **when and where an event occurred**, in terms of the process that executed it and the location in the source code.



# Generating Event Traces

- To enable tracing with TAU, set the environment variable **TAU\_TRACE** to 1.
- Similarly you can enable/disable profile with the **TAU\_PROFILE** variable.
- Set the output directory with a environment variable:

```
% setenv TRACEDIR /users/sameer/tracedata
```



# Content

- Overview
- Instrumentation
- Profiling
- Tracing
- Analyzing
- Application Scenario



# Analyzing -- Text summary

- For a quick view summary of TAU performance, use

```
%> cd MULTI__P_WALL_CLOCK_TIME  
%> pprof
```

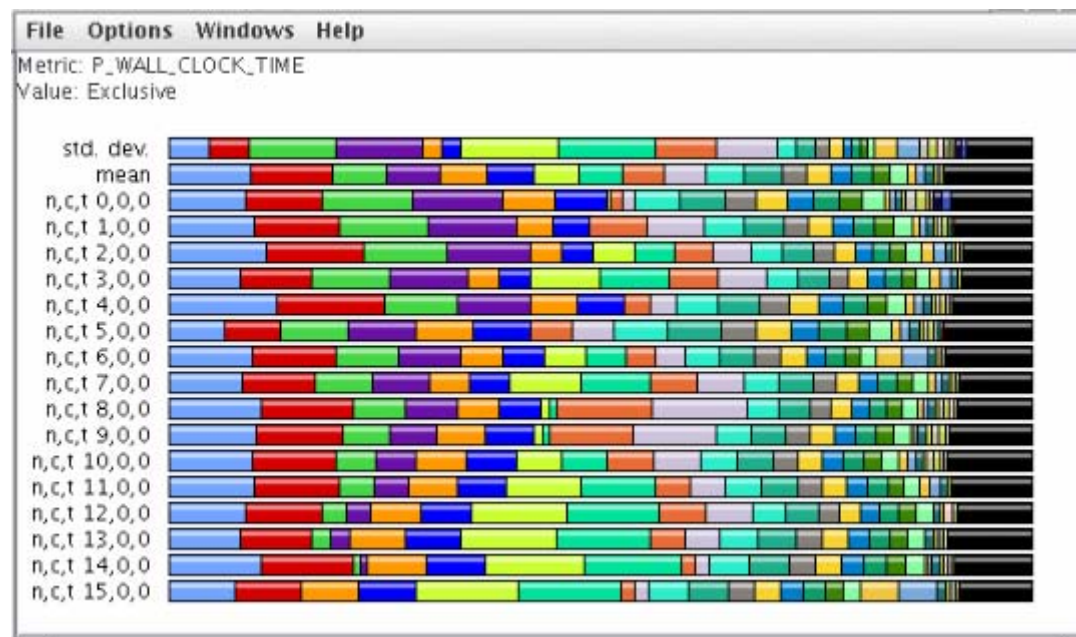
- It reads and prints a summary of the TAU data in the current directory.

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	24	590	1	1	590963	main
95.9	26	566	1	2	566911	multiply
47.3	279	279	1	0	279280	multiply-opt
44.1	260	260	1	0	260860	multiply-regul



# Analyzing -- ParaProf

- To launch ParaProf, execute paraprof from the command line where the profiles are located.






# Analyzing -- Jumpshot

- To use Argonne's Jumpshot (bundled with TAU), first merge and convert TAU traces to slog2 format:

```
% tau_treemerge.pl  
% tau2slog2 tau.trc tau.edf -o tau.slog2  
% jumpshot tau.slog2
```





# Content


- Overview
- Instrumentation
- Profiling
- Tracing
- Analyzing
- [Application Scenario](#)



# Application Scenario (1)

- Q. What routines account for the most time? How much?
- A. Create a flat profile with wallclock time.

Metric: P\_VIRTUAL\_TIME  
Value: Exclusive  
Units: seconds

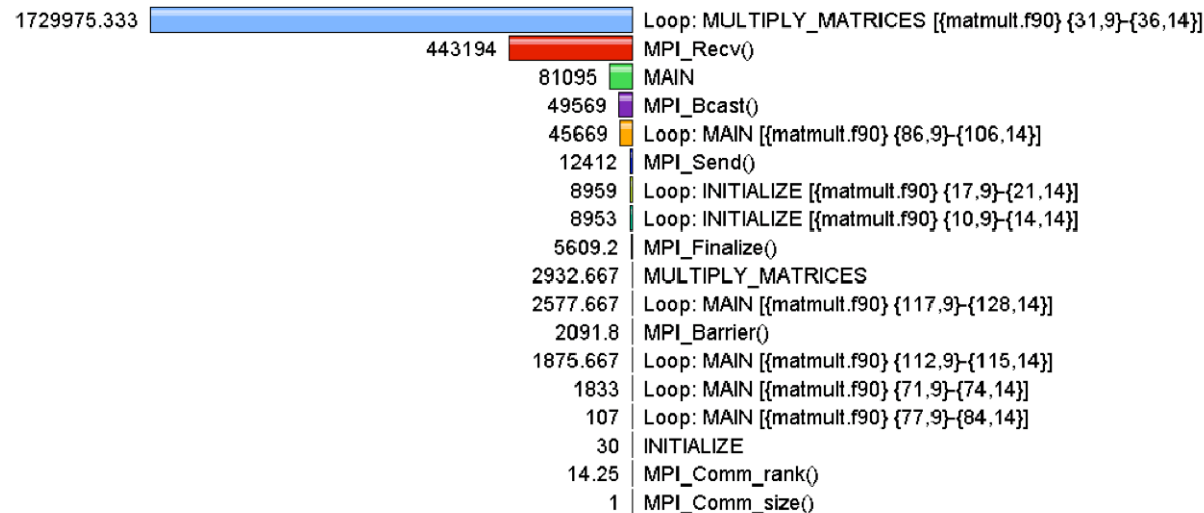




# Application Scenario (2)

- Q. What loops account for the most time? How much?
- A. Create a flat profile with loop instrumentation.

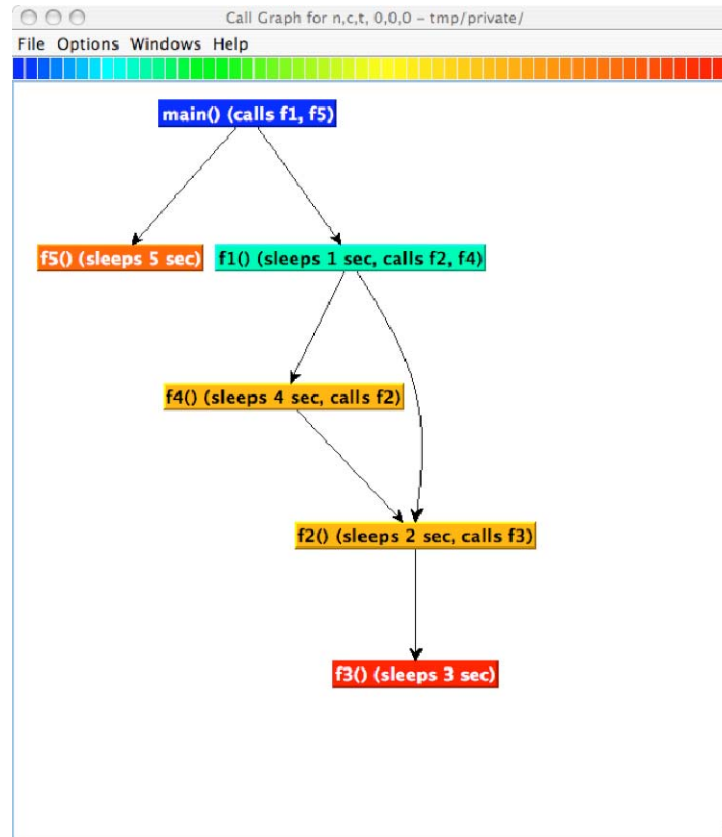
Metric: GET\_TIME\_OF\_DAY  
Value: Exclusive  
Units: microseconds





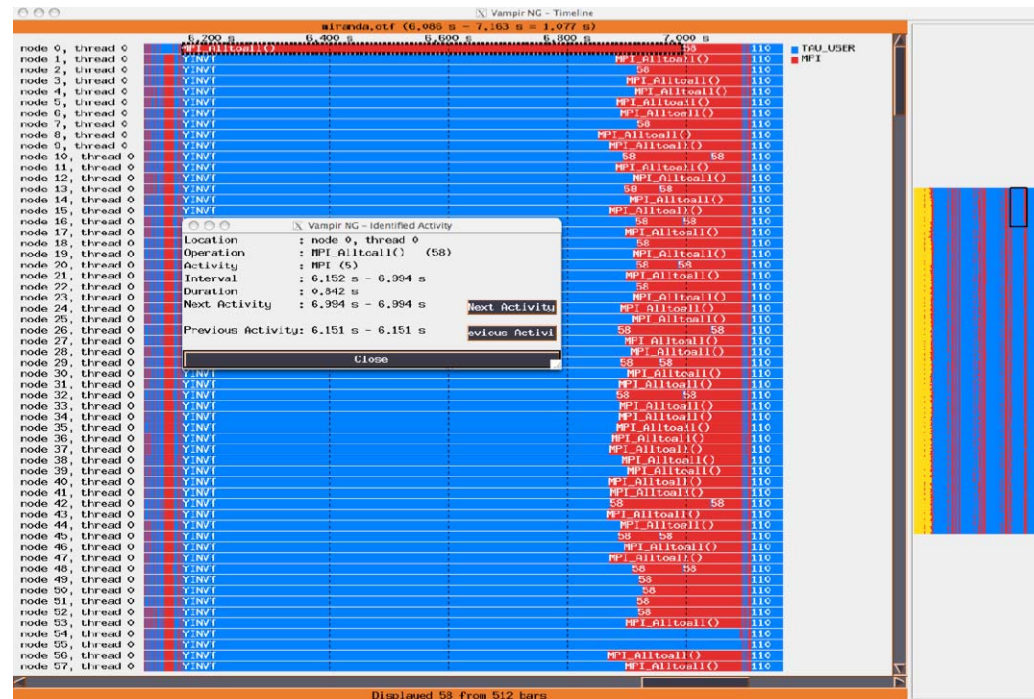
# Application Scenario (3)

- Q. Who calls MPI\_Barrier() Where?
- A. Create a callpath profile with given depth.



# Application Scenario (4)

- Q. What happens in my code at a given time ?
- A. Create an event tree.





# Application Scenarios

- Q. How do I instrument Python Code?
- A. Create an python wrapper library.
- Q. How does my application scale?
- A. Examine profiles in PerfExplorer.
- Q. What MFlops am I getting in all loops?
- A. Create a flat profile with PAPI\_FP\_INS/OPS and time with loop instrumentation.





# Using TAU: Simplest Case

## ■ Uninstrumented code:

- % mpirun -np 8 ./a.out

## ■ With TAU:

- % mpirun -np 8 **tau\_exec** ./a.out

- % paraprof



# references

## ■ TAU User Guide

□ <https://www.cs.uoregon.edu/research/tau/docs/newguide/bk01pt01.html>

## ■ TAU Website:

□ <http://tau.uoregon.edu>

## ■ Sameer Shende, *TAU Performance System*, Performance Research Lab, University of Oregon.

□ <http://tau.uoregon.edu/tau.ppt>